



Danny Hermelin

Fairness in Repetitive Scheduling

Joint work with Klaus Heeger, Yuval Itzhaki, George B. Mertzios, Hendrik Molter, Rolf Niedermeier, Michael L. Pinedo, Danny Segev, and [Dvir Shabtay](#)

schedulingseminar.com

Classical Scheduling

- ▶ Most single machine scheduling problems look like this:

- ▶ There are n jobs (or clients):



- ▶ Each job j has:

- ▶ a processing time p_j ,
 - ▶ a weight w_j (which may equal 1),
 - ▶ a due date d_j (sometimes),
 - ▶ a release time r_j (sometimes),
 - ▶ etc...

Classical Scheduling

- ▶ Most single machine scheduling problems look like this:
 - ▶ A **schedule** (when there are no release times) is simply a permutation of the jobs, specifying the order of processing:



- ▶ The **completion time** of job j in such a schedule is $C_j = \sum_{\substack{i \text{ is not after } j \\ \text{in the schedule}}} p_i$

Classical Scheduling

- ▶ The clients may have different objectives:
 - ▶ they may want to -



- ▶ minimize **completion time**,

$$C_j = \sum_{\pi(i) \leq \pi(j)} p_i$$


- ▶ minimize **lateness**,

$$L_j = C_j - d_j$$

- ▶ not be tardy,

$$U_j = \begin{cases} 1 & : C_j > d_j \\ 0 & : C_j \leq d_j \end{cases}$$

- ▶ etc...




Classical Scheduling

- ▶ The scheduler, being the service provider, typically decides the objective of the schedule.
- ▶ The scheduler might decide to try to be as fair as possible to all clients.
- ▶ One way to do so, is to minimize their total completion time - $\sum_j c_j$
 - ▶ which is equivalent to minimizing the average completion time of a client.



Classical Scheduling


- ▶ To minimize $\sum C_j$, use the **Shortest Processing Time first (SPT)** rule:



$$\sum_j C_j = 3p_1 + 2p_2 + p_3$$

Classical Scheduling


- ▶ Suppose the clients just don't want to be late.
- ▶ To minimize $\sum U_j$, use the **Earliest Deadline First (EDD)** rule:



- ▶ Scan the jobs in EDD order, and when first encountering a tardy job, remove the largest job already processed (Hodson-Moore Alg.).
- ▶ Is this really fair?

Repetitive Scheduling

- But what happens when the clients keep on returning ?



- Now, in the simplest case, each client has m jobs, one per each of the m days of the month.
- Let $p_{i,j}$ denote the processing time of client j 's job on day i .
- Define $C_{i,j}$ to be the completion time of client j 's job on day i .
- Similarly define $w_{i,j}$, $d_{i,j}$, and $r_{i,j}$ when necessary.

Repetitive Scheduling

- ▶ If we minimize the total completion time $\sum C_{i,j}$, a single client can complete last on every day!



- ▶ This is particularly unfair if the job processing time aren't entirely determined by the clients.
- ▶ For example:
 - ▶ Healthcare.
 - ▶ Civic duties (jury duty, military service, ect..).
 - ▶ Etc...



Repetitive Scheduling

- ▶ If we minimize the total completion time $\sum U_{i,j}$, a single client can be tardy every day!



- ▶ This means that from his point of view, he never gets any service!
- ▶ This is particularly unfair when the job due dates aren't entirely determined by the clients.



Fair Repetitive Scheduling

- ▶ Instead, it makes much more sense to minimize

$$\max_{j \in [n]} C_j = \max_{j \in [n]} \sum_{i \in [m]} C_{i,j}$$

in case the clients are interested in minimizing their completion time.



- ▶ Or, to minimize

$$\max_{j \in [n]} U_j = \max_{j \in [n]} \sum_{i \in [m]} U_{i,j}$$

in case the clients are interested in not being tardy.

Fair Repetitive Scheduling

- ▶ More generally, we may have any (say, minimization) objective function $F_{i,j} = f(C_{i,j})$ that typically depends on the completion time of the clients on any given day.
- ▶ We define the (single machine) **Fairness in Repetitive Scheduling** problem, $1|\text{rep}|\max_j \sum_i F_{i,j}$, as the problem of minimizing

$$\max_{j \in [n]} F_j = \max_{j \in [n]} \sum_{i \in [m]} F_{i,j}$$



- ▶ In the **decision version**, we are given a **fairness threshold** K , and the goal is to determine whether there exists a schedule with $\max F_j \leq K$.

Fair Repetitive Scheduling

- ▶ Thus, we obtain a **performance matrix**, which specifies the performance (according to the given objective function) of each job.

$F_{1,1}$	$F_{1,2}$	$F_{1,n}$
$F_{2,1}$	$F_{2,2}$	$F_{2,n}$
...
...
...
$F_{m,1}$	$F_{m,2}$	$F_{m,n}$
F_1	F_2	F_n


rows
correspond
to days

column j
sums up to F_j

- ▶ The goal is to minimize the maximum value in the final row.

Fair Repetitive Scheduling

- ▶ Our model is robust, and can handle several tweaks and changes, such as




- ▶ More complex machine models.
 - ▶ E.g., the $Pq|rep|\max_j \sum_i C_{i,j}$ problem is the problem of minimizing the maximum total completion time of every client when there are q parallel machines every day.
- ▶ Different fairness thresholds K_j for each client.
 - ▶ E.g., in cases where there are premium customers, etc.



Fair Repetitive Scheduling

- ▶ Our model allows several interesting variants to study -



- ▶ Is there a schedule which is k -fair for all but s clients?

- ▶ A variant of **approximate fairness** where a solution is fair to almost everyone.



- ▶ The price of fairness?
 - ▶ The ratio between the global optimum and the worst k -fair solution.

The $1 | \text{rep} | \max \sum C_{i,j}$ Problem

- ▶ Consider first parameter n = number of clients.

Theorem: *The $1 | \text{rep} | \max \sum C_{i,j}$ problem is (weakly) NP-hard for two or more clients ($n \geq 2$).*

Theorem: *The $1 | \text{rep} | \max \sum C_{i,j}$ problem is pseudo polynomial-time solvable for a constant number of clients ($n = O(1)$).*

- ▶ For parameter k = fairness threshold, we can show -

Theorem: *The $1 | \text{rep} | \max \sum C_{i,j}$ problem is (strongly) NP-hard even for constant fairness thresholds ($k \geq 37$).*

- ▶ Note that this shows that the problem is APX-hard, meaning it doesn't admit a PTAS (most likely).

The $1 | \text{rep} | \max \sum C_{i,j}$ Problem

- ▶ Next consider parameter m = number of days. We prove

Theorem: *The $1 | \text{rep} | \max \sum C_{i,j}$ problem is (weakly) NP-hard for four or more days ($m \geq 4$).*

Theorem: *The $1 | \text{rep} | \max \sum C_{i,j}$ problem is polynomial-time solvable for two days ($m = 2$).*

- ▶ This leads to the first open problem of the talk:



**What is the complexity
of $1 | \text{rep} | \max \sum C_{i,j}$ for
three days??**

The Two Days Algorithm

- ▶ The algorithm is inspired by Johnson's algorithm for $F2||C_{\max}$.
- ▶ We begin with a structural lemma:

Lemma (Property 1): *There is an optimal schedule for any $1|rep|\max \sum C_{i,j}$ instance on two days, where the schedule on the second day is in reverse order of the schedule on the first day.*

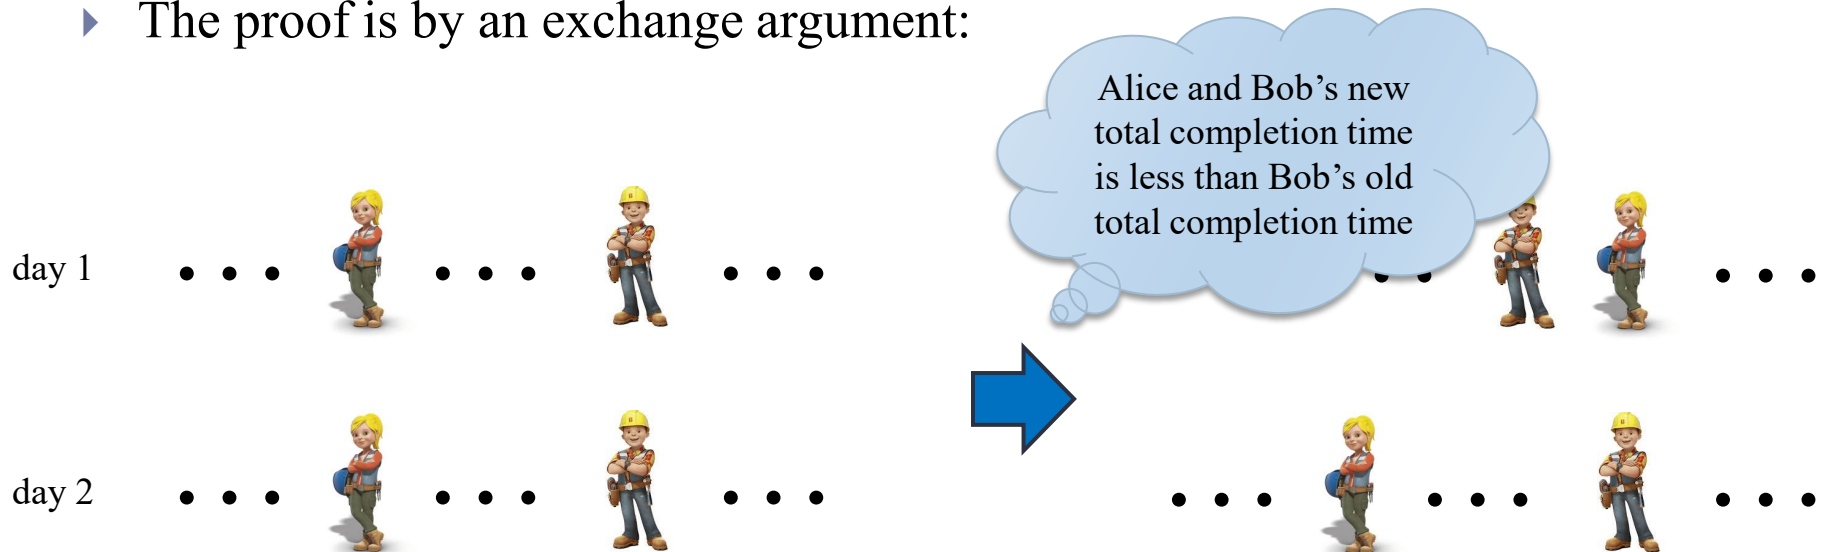


- ▶ Thus, it is enough to only determine the order of processing on the first day.

The Two Days Algorithm

Lemma (Property 1): *There is an optimal schedule for any $1|\text{rep}|\max \sum C_{i,j}$ instance on two days, where the schedule on the second day is in reverse order of the schedule on the first day.*

- The proof is by an exchange argument:



The Two Days Algorithm

- ▶ Using the lemma, we construct an optimal schedule as follows:
- ▶ Partition the clients into two types:

type 1



smaller processing
time on day 1 at least
as large as on day 2

$$p_{1,j} \leq p_{2,j}$$



smaller processing
time on day 2

$$p_{1,j} > p_{2,j}$$

type 2


- ▶ On **day 1**: Schedule first the jobs of type 1 clients in SPT order, and then the jobs of type 2 clients in LPT order.
- ▶ On **day 2** do the reverse of day 1.

The Two Days Algorithm

Lemma: *The schedule constructed is optimal.*

- ▶ **Proof:** Consider only optimal solutions where property 1 holds.
 - ▶ **Step 1:** Show that this set includes solutions where type 1 clients are scheduled before type 2 clients on day 1 (Property 2).

If not there always exists a type 1 job scheduled directly after a type 2 job




The Two Days Algorithm

Lemma: *The schedule constructed is optimal.*

- ▶ **Proof:** Consider only optimal solutions where property 1 and 2 hold.
 - ▶ **Step 2:** Show that this set only includes solutions where on day 1, type 1 clients are scheduled in SPT order, and type 2 clients scheduled in LPT order (Property 3).

If not, there always is a consecutive pair of clients which can be swapped.



Parameterized Complexity

- ▶ We can also show the following -

Theorem: *The $1|\text{rep}|\max \sum C_{i,j}$ problem is strongly $W[1]$ -hard when parameterized by the number of days.*

- ▶ This implies that the problem is unlikely to admit a $f(m)n^{O(1)}$ algorithm, even if the processing times are given in unary.
 - ▶ A $n^{f(m)}$ algorithm is still possible (yet unknown) in this case.
- ▶ On the other hand, we can show that -

Theorem: *The $1|\text{rep}|\max \sum C_{i,j}$ problem can be solved in $f(K+m)n^{O(1)}$ time.*

- ▶ The proof uses **n-fold** ILPs.

Approximation Algorithms

- ▶ We say that a solution is α -approximate if it has a fairness threshold of k such that $k \leq \alpha \cdot OPT$, where OPT is the optimal fairness threshold.
- ▶ We show that -

Theorem: *The $1|rep|\max \sum C_{i,j}$ problem admits a 2-approximation algorithm that runs in polynomial-time.*

- ▶ While the algorithm runs in polynomial-time, it relies on several applications of an LP solver.



**Is there a more efficient
(combinatorial)
constant approximation
algorithm??**

2-Approximation Algorithm

- ▶ The algorithm is inspired by 3-approximation algorithm for $1|r_j|\sum w_j C_j$ of Hall, Schulz, D. B. Shmoys, and Wein [MathofOR'97].
- ▶ Consider the following LP:

$$\begin{array}{ll}\min & K \\ \text{s.t.} & \sum_{i \in [m]} x_{i,j} \leq K \quad \forall j \in [n] \\ & \sum_{j \in S} p_{i,j} x_{i,j} \geq \frac{1}{2} \cdot P_i^2(S) \quad \forall i \in [m], S \subseteq [n]\end{array}$$

- ▶ The variables are:
 - ▶ $x_{i,j}$ = completion time of client j 's job on day i .
 - ▶ K = the fairness threshold.

2-Approximation Algorithm

- ▶ The first set of constraints ensures that no client has total completion time which exceeds the fairness threshold k :

$$\sum_{i \in [m]} x_{i,j} \leq K \quad \forall j \in [n]$$

- ▶ The second set is less clear -

$$\sum_{j \in S} p_{i,j} x_{i,j} \geq \frac{1}{2} \cdot P_i^2(S) \quad \forall i \in [m], S \subseteq [n]$$

- ▶ Here $P_i^2(S)$ denotes the total processing time of the jobs of S on day i , squared.

2-Approximation Algorithm

- ▶ The constraint is clearly satisfied for every singleton $S=\{j\}$ since

$$p_{i,j} \cdot C_{i,j} \geq p_{i,j} \cdot p_{i,j} > \frac{1}{2} \cdot p_{i,j}^2 = \frac{1}{2} \cdot P_i^2(S).$$

- ▶ And for $|S|>1$ we have

$$\begin{aligned} \sum_{j \in S} p_{i,j} C_{i,j} &\geq \sum_{j,k \in S, j \leq k} p_{i,j} p_{i,k} = \frac{1}{2} \cdot \left(\sum_{j \in S} p_{i,j} \right)^2 + \frac{1}{2} \sum_{j \in S} p_{i,j}^2 \\ &> \frac{1}{2} \cdot \left(\sum_{j \in S} p_{i,j} \right)^2 = \frac{1}{2} \cdot P_i^2(S). \end{aligned}$$

- ▶ As either the job of client j is scheduled before the job of client k , or vice-versa.
- ▶ Thus, every schedule satisfies the second set of constraints.

2-Approximation Algorithm

$$\begin{array}{ll} \min & K \\ \text{s.t.} & \sum_{i \in [m]} x_{i,j} \leq K \quad \forall j \in [n] \\ & \sum_{j \in S} p_{i,j} x_{i,j} \geq \frac{1}{2} \cdot P_i^2(S) \quad \forall i \in [m], S \subseteq [n] \end{array}$$

- ▶ Note that our LP has an **exponential number of constraints**, one for each subset of clients.
- ▶ Lucky, the LP has a **separation oracle** (due to Queyranne [MathProg'93]) -
 - ▶ A separation oracle is a **polynomial-time algorithm** that receives a solution x to the LP and determines **whether x is feasible or not**. If x is not feasible, the oracle returns a constraint which is violated by x .
- ▶ One can use the Primal-Dual method, in conjunction with the oracle, to obtain an optimal solution for the LP.

2-Approximation Algorithm

- ▶ Let $x_{1,1}^*, \dots, x_{n,m}^*$ denote an optimal solution for the LP.
- ▶ Using this solution, we construct a schedule for each day i by **processing the jobs in non-decreasing values of $x_{i,j}^*$** .

schedule π_i
for day i



$x_{i,j}^*$



$x_{i,j}^*$




$x_{i,j}^*$

- ▶ The completion time of client j 's job on day i is $C_{i,j} = \sum_{\pi_i(k) \leq \pi_i(j)} p_{i,k}$.

2-Approximation Algorithm

Lemma: *The schedule constructed is 2-approximate.*

- ▶ **Proof:** Fix $i \in [m]$ and $j \in [n]$, and consider the job of client j on day i .
- ▶ Let S_ℓ denote the set of clients whose jobs were not scheduled after the job of client j (i.e., j and those before him on day i).



- ▶ Then, by our construction, we have $C_{i,j} = \sum_{\pi_i(k) \leq \pi_i(j)} p_{i,k} = P_i(S_\ell)$.

2-Approximation Algorithm

Lemma: *The schedule constructed is 2-approximate.*

- ▶ **Proof:** Fix $i \in [m]$ and $j \in [n]$, and consider the job of client j on day i .
- ▶ Then, by our construction, we have $C_{i,j} = \sum_{\pi_i(k) \leq \pi_i(j)} p_{i,k} = P_i(S_\ell)$.
- ▶ On the other hand, as $x_{1,1}^*, \dots, x_{n,m}^*$ is feasible, and $x_{i,j}^* \geq x_{i,k}^*$ for all clients $k \in S_\ell$, we have

$$\frac{1}{2} \cdot P_i^2(S_\ell) \leq \sum_{k \in S_\ell} p_{i,k} x_{i,k}^* \leq x_{i,j}^* \cdot \sum_{k \in S_\ell} p_{i,k} = x_{i,j}^* \cdot P_i(S_\ell).$$

- ▶ It follows that $2x_{i,j}^* \geq P_i(S_\ell) = C_{i,j}$
- ▶ Thus, for any client $j \in [n]$, we get $\sum_{i=1}^n C_{i,j} \leq 2 \sum_{i=1}^n x_{i,j}^* \leq 2K^*$,

where K^* is the optimal fairness threshold.

Further Results

- ▶ We can also show the following -

Theorem: *The $1|rep|\max \Sigma C_{i,j}$ problem admits a PTAS for a constant number of days.*

- ▶ The PTAS relies on an involved **batching scheme** where we batch groups of jobs together.
- ▶ We also consider the **all days are the same** case, and show -

Theorem: *The $1|rep, p_{i,j}=p_j|\max \Sigma C_{i,j}$ problem admits a $\frac{(1+\sqrt{2})}{2}$ -approximation algorithm that runs in near linear-time.*

Theorem: *The $1|rep, p_{i,j}=p_j|\max \Sigma C_{i,j}$ problem admits a QPTAS.*

The $1 | \text{rep} | \max \sum U_{i,j}$ Problem

- ▶ The problem is already quite hard, even when each client has the same due date on each day:

Theorem: *The $1 | \text{rep}, d_{i,j}=d_j | \max \sum U_{i,j}$ problem is NP-hard, even for $K=1$.*

- ▶ In the all days are the same variant, we can show for the single due date case a 1.5-approximation algorithm
- ▶ On the other hand, it is easy for unit processing times:

Theorem: *The $1 | \text{rep}, p_{i,j}=1 | \max \sum U_{i,j}$ problem is polynomial-time solvable.*

- ▶ This theorem holds even when the jobs have release times, and there are multiple parallel machines for processing.

The $1 | \text{rep} | \max \sum Z_{i,j}$ Problem

- ▶ This problem is quite hard as well:

Theorem: *The $1 | \text{rep} | \max \sum Z_{i,j}$ problem is polynomial-time solvable for $k \in \{m - 1, m\}$, and NP-hard otherwise.*

- ▶ However, when the number of clients is small -

Theorem: *The $1 | \text{rep} | \max \sum Z_{i,j}$ problem can be solved in $f(n)m^{O(1)}$ time.*

- ▶ And when then the number of days is small -

Theorem: *The $1 | \text{rep}, d_i, j=d_j | \max \sum Z_{i,j}$ problem can be solved in $O(mn)^{(m+1)}$ time.*

- ▶ The **all days are the same** case translates directly to **interval graph coloring**.

Bibliography

1. *Klaus Heeger, **D.H.**, George B. Mertzios, Hendrik Molter, Rolf Niedermeier, Dvir Shabtay: Equitable scheduling on a single machine. J. Sched. 26(2): 209-225 (2023).*
2. ***D.H.**, Hendrik Molter, Rolf Niedermeier, Michael L. Pinedo, Dvir Shabtay: Fairness in repetitive scheduling. Eur. J. Oper. Res. 323(3): 724-738 (2025).*
3. *Klaus Heeger, **D.H.**, Yuval Itzhaki, Hendrik Molter, Dvir Shabtay: Fair repetitive interval scheduling. Algorithmica 87(9): 1340-1368 (2025).*
4. ***D.H.**, Danny Segev, Dvir Shabtay: Approximate fair repetitive scheduling. Work in progress.*

THANK YOU
FOR YOUR
ATTENTION